

# GenAI Assignment Documentation

## Multi-Agent AI Ops Assistant

**Candidate:** Manas Kumar Rout

**Company:** TrulyMadly

**Role:** GenAI Intern

---

### 1. Project Overview

This project implements a **Multi-Agent AI Operations Assistant** that can understand a user task, plan execution steps, call real third-party APIs, validate results, and return a complete final response.

The system is designed to demonstrate:

- Multi-agent reasoning
- Structured LLM outputs
- Real API integrations
- End-to-end execution
- Local execution with a single command

The project satisfies all **mandatory pass/fail criteria** specified in the assignment email.

---

### 2. High-Level Architecture

#### Architecture Flow

User Input



Planner Agent (LLM → Structured JSON Plan)



Executor Agent (Tool/API Calls)



Verifier Agent (Validation + Retry)



Final Answer

## Key Design Principles

- Separation of concerns via agents
  - Deterministic structured planning
  - Tool-based execution (no hardcoded responses)
  - Fault tolerance via verifier retry
- 

## 3. Multi-Agent Design

### 3.1 Planner Agent

#### Purpose:

Transforms a natural-language task into a **structured JSON execution plan**.

#### Responsibilities:

- Understand user intent
- Decide which tools/APIs are required
- Output a machine-readable plan

#### Example Output:

```
{  
  "task": "Get current weather in Mumbai",  
  "steps": [  
    {  
      "tool": "weather.current_by_city",  
      "tool_input": { "city": "Mumbai", "units": "metric" },  
      "save_as": "mumbai_weather"  
    }  
  ]  
}
```

---

### **3.2 Executor Agent**

**Purpose:**

Executes each step from the planner by invoking real APIs.

**Responsibilities:**

- Call registered tools
  - Handle API responses
  - Store outputs in a structured format
- 

### **3.3 Verifier Agent**

**Purpose:**

Ensures the task has been fully satisfied.

**Responsibilities:**

- Validate tool outputs against the task
  - Detect missing or incomplete results
  - Retry missing tool calls **once**
  - Produce final user-ready output
- 

## **4. Integrated Third-Party APIs**

### **4.1 Groq LLM API**

- Used for Planner agent
- Model: llama-3.3-70b-versatile
- Provides reasoning and structured output

### **4.2 OpenWeatherMap API**

- Fetches real-time weather data
- Endpoint: /data/2.5/weather
- Used by weather.current\_by\_city tool

### 4.3 GitHub Search API

- Searches public GitHub repositories
  - Used by `github.search_repos` tool
  - Token optional (avoids rate limits)
- 

### 5. Project Structure

```
trulymadly-genai-assignment/
```

```
|  
|   └── agents/  
|       |   └── planner.py  
|       |   └── executor.py  
|       |   └── verifier.py  
|       |   └── schemas.py  
|  
|  
|   └── tools/  
|       |   └── github_tool.py  
|       |   └── weather_tool.py  
|       |   └── tool_registry.py  
|  
|  
└── llm/  
    |   └── client.py  
    |   └── prompts.py  
|  
|  
└── main.py  
└── requirements.txt  
└── .env.example
```

|— README.md

---

## 6. Environment Configuration

### Required Environment Variables

.env (not committed):

OPENAI\_API\_KEY=gsk\_XXXXXXXXXX

OPENAI\_BASE\_URL=https://api.groq.com/openai/v1

OPENAI\_MODEL=llama-3.3-70b-versatile

OPENWEATHER\_API\_KEY=XXXXXXXXXX

GITHUB\_TOKEN=XXXXXXXXXX

### .env.example (committed):

OPENAI\_API\_KEY=

OPENAI\_BASE\_URL=https://api.groq.com/openai/v1

OPENAI\_MODEL=llama-3.3-70b-versatile

OPENWEATHER\_API\_KEY=

GITHUB\_TOKEN=

---

## 7. Setup Instructions (Localhost)

### 7.1 Prerequisites

- Python 3.10+
- Internet connection

### 7.2 Installation

python -m venv .venv

.venv\Scripts\activate # Windows

```
pip install -r requirements.txt
```

---

## 8. Running the Project

The project runs locally using **one command**, as required:

```
uvicorn main:app --reload
```

Open in browser:

```
http://localhost:8000
```

---

## 9. Example Prompts for Evaluation

### 1. Weather Query

What's the weather in Mumbai right now?

### 2. GitHub Repository Search

Find top 5 GitHub repos for "fastapi jwt auth" and summarize best choice.

### 3. Multi-Tool Query

Find top 5 repos for "react markdown editor" and also check weather in Bangalore.

### 4. Another Weather Query

What's the current weather in Delhi?

---

## 10. Sample Successful Output

```
{
```

```
  "task": "What's the weather in Mumbai right now?",
```

```
  "plan": {...},
```

```
  "tool_outputs": {...},
```

```
  "final_answer": {
```

```
    "city": "Mumbai",
```

```
    "weather": "smoke",
```

```
"temperature": 30.99,  
"humidity": 40  
,  
"verifier_notes": [  
    "Current weather in Mumbai is provided as requested."  
]  
}
```

---

## 11. Known Limitations & Trade-offs

- Single-turn execution (no conversational memory)
  - Verifier retries only once (prevents infinite loops)
  - No user authentication (out of scope)
  - Minimal UI (focus on backend GenAI logic)
- 

## 12. Security Considerations

- .env is excluded from GitHub via .gitignore
  - No API keys are hardcoded
  - Push protection enabled by GitHub
  - .env.example contains placeholders only
- 

## 13. Assignment Requirements Checklist

Requirement	Status
Multi-agent design	✓
Structured LLM output	✓
≥2 real APIs	✓

Requirement	Status
End-to-end execution	✓
One-command local run	✓
No hardcoded responses	✓

---

## 14. Conclusion

This project demonstrates a complete **GenAI multi-agent system** with real-world API integrations, structured reasoning, fault tolerance, and evaluator-friendly execution.

The system fulfills all requirements outlined by TrulyMadly and is ready for technical evaluation.

---

### GitHub Repository:

<https://github.com/Manas1024/Trulymadly-genai-assignment>